

Part 2: Indexing Analysis

Query 1: Average Price of Watches per Country

SQL Query

```
SELECT wb.hq_country, ROUND(AVG(w.retail_price), 2) AS avg_price
FROM Watches w
JOIN Watch_Brands wb ON w.brand_name = wb.brand_name
GROUP BY wb.hq_country
ORDER BY avg_price DESC;
```

Indexes Tried

```
CREATE INDEX idx_brand_name_watches ON Watches(brand_name);
CREATE INDEX idx_watch_brands_country ON Watch_Brands(hq_country);
CREATE INDEX idx_brand_price ON Watches(brand_name, retail_price);
```

Summary Table

Index Strategy	Cost Estimate	Actual Time (ms)
No indexes	86893	196–286
Join indexes on brand_name	4887	1.11–16.2
+ Index on hq_country	4887	1.00–15.6
+ Composite index on brand_name, price	4887	1.07–15.4

```

mysql> EXPLAIN ANALYZE
-> SELECT wb.hq_country, ROUND(AVG(w.retail_price), 2) AS avg_price
-> FROM Watches w
-> IGNORE INDEX (PRIMARY, idx_retail_price_brand, idx_retail_price_brand, idx_brand_price, idx_brand_name_watches, idx_case_material, idx_brand_case) -- Ignore indexes on the Watches table
-> JOIN Watch_Brands wb
-> IGNORE INDEX (PRIMARY, idx_hq_country, idx_watch_brands_country, idx_brand_name_watch_brands, idx_group_brand) -- Ignore indexes on the Watch_Brands table
-> ON w.brand_name = wb.brand_name
-> GROUP BY wb.hq_country
-> ORDER BY avg_price DESC;
+-----+
| EXPLAIN
+-----+
|
+-----+
| -> Sort: avg_price DESC (actual time=38.6..38.6 rows=5 loops=1)
-> Table scan on <temporary> (actual time=37.4..37.4 rows=5 loops=1)
-> Aggregate using temporary table (actual time=37.4..37.4 rows=5 loops=1)
-> Inner hash join (w.brand_name = wb.brand_name) (cost=86983 rows=20214) (actual time=2.86..25.2 rows=17884 loops=1)
-> Table scan on w (cost=1.96 rows=17384) (actual time=0.332..1.7 rows=17884 loops=1)
-> Hash
-> Table scan on wb (cost=5.25 rows=50) (actual time=0.0525..0.0832 rows=50 loops=1)
+-----+
|
+-----+
1 row in set (0.07 sec)

mysql> explain analyze SELECT wb.hq_country, ROUND(AVG(w.retail_price), 2) AS avg_price
-> FROM Watches w
-> JOIN Watch_Brands wb ON w.brand_name = wb.brand_name
-> GROUP BY wb.hq_country
-> ORDER BY avg_price DESC;
+-----+
| EXPLAIN
+-----+
|
+-----+
| -> Sort: avg_price DESC (actual time=19.5..19.5 rows=5 loops=1)
-> Stream results (cost=4887 rows=6) (actual time=1.57..19.5 rows=5 loops=1)
-> Group aggregate: avg(w.retail_price) (cost=4887 rows=6) (actual time=1.57..19.5 rows=5 loops=1)
-> Nested loop inner join (cost=2866 rows=20214) (actual time=0.369..14.4 rows=17884 loops=1)
-> Covering index scan on wb using idx_hq_country (cost=5.25 rows=50) (actual time=0.112..0.152 rows=50 loops=1)
-> Covering index lookup on w using idx_retail_price_brand (brand_name=wb.brand_name) (cost=17.6 rows=404) (actual time=0.0486..0.261 rows=358 loops=50)
+-----+
|
+-----+
1 row in set (0.04 sec)

```

Analysis

Overall the cost estimate decreased significantly from the join index on brand_name while remaining constant on the additional indexes. This shows a strong positive index and we must take this into consideration in the final analysis.

Query 2: Average Price per Brand

SQL Query

```
SELECT wb.brand_name, ROUND(AVG(w.retail_price), 2) AS avg_price  
FROM Watches w  
JOIN Watch_Brands wb ON w.brand_name = wb.brand_name  
GROUP BY wb.brand_name  
ORDER BY avg_price DESC;
```

Indexes Tried

```
CREATE INDEX idx_brand_name_watch_brands ON Watch_Brands(brand_name);  
CREATE INDEX idx_retail_price_brand ON Watches(brand_name, retail_price);
```

Summary Table

Index Strategy	Cost Estimate	Actual Time (ms)
No indexes	86981	~20.2
Index on brand_name (Watch_Brands)	4887	~13.8
Composite index (brand + price)	4887	~14.9

```
mysql> EXPLAIN ANALYZE SELECT wb.brand_name, ROUND(AVG(w.retail_price), 2) AS avg_price FROM Watches w IGNORE INDEX (PRIMARY, idx_retail_price_brand, idx_retail_price_brand,
ch_Brands wb IGNORE INDEX (PRIMARY, idx_hq_country, idx_watch_brands_country, idx_brand_name_watch_brands, idx_group_brand) ON w.brand_name = wb.brand_name GROUP BY wb.brand
+-----+
| EXPLAIN
+-----+
|
+-----+
| -> Sort: avg_price DESC (actual time=22.8..22.8 rows=43 loops=1)
  -> Table scan on <temporary> (actual time=22.7..22.7 rows=43 loops=1)
    -> Aggregate using temporary table (actual time=22.7..22.7 rows=43 loops=1)
      -> Inner hash join (w.brand_name = wb.brand_name) (cost=86981 rows=20214) (actual time=0.141..12.5 rows=17884 loops=1)
        -> Table scan on w (cost=1.91 rows=17384) (actual time=0.0417..7.54 rows=17884 loops=1)
        -> Hash
          -> Table scan on wb (cost=5.25 rows=50) (actual time=0.0527..0.0631 rows=50 loops=1)
+-----+
1 row in set (0.03 sec)

mysql> explain analyze SELECT wb.brand_name, ROUND(AVG(w.retail_price), 2) AS avg_price
-> FROM Watches w
-> JOIN Watch_Brands wb ON w.brand_name = wb.brand_name
-> GROUP BY wb.brand_name
-> ORDER BY avg_price DESC;
+-----+
| EXPLAIN
+-----+
|
+-----+
| -> Sort: avg_price DESC (actual time=16.8..16.8 rows=43 loops=1)
  -> Stream results (cost=4887 rows=50) (actual time=0.315..16.7 rows=43 loops=1)
    -> Group aggregate: avg(w.retail_price) (cost=4887 rows=50) (actual time=0.313..16.7 rows=43 loops=1)
      -> Nested loop inner join (cost=2866 rows=20214) (actual time=0.112..11.9 rows=17884 loops=1)
        -> Covering index scan on wb using idx_brand_name_watch_brands (cost=5.25 rows=50) (actual time=0.0453..0.0813 rows=50 loops=1)
        -> Covering index lookup on w using idx_retail_price_brand (brand_name=wb.brand_name) (cost=17.6 rows=404) (actual time=0.0287..0.21 rows=358 loops=50)
+-----+
1 row in set (0.02 sec)
```

Analysis

There was once again a very large difference just like query 1.

Query 3: Watches That Exceed Market Price

SQL Query

```
SELECT wb.brand_name, ROUND(AVG(w.retail_price), 2) AS avg_retail_price
FROM Watches w
JOIN Watch_Brands wb ON w.brand_name = wb.brand_name
GROUP BY wb.brand_name
HAVING AVG(w.retail_price) > (
    SELECT AVG bmp.Price
    FROM Brand_Market_Data bmp
    WHERE bmp.brand_name = wb.brand_name
);
```

Indexes Tried

```
CREATE INDEX idx_brand_name_bmd ON Brand_Market_Data(brand_name);
CREATE INDEX idx_bmd_brand_price ON Brand_Market_Data(brand_name, price);
CREATE INDEX idx_brand_price ON Watches(brand_name, retail_price);
```

Summary Table

Index Strategy	Cost Estimate	Actual Time (ms)
No indexes	88715	~280
Index on brand_name (BMD)	4887	~54.1
Composite index on (brand_name, price)	4887	~52–54
Composite index on Watches	4887	~50–53

Analysis

```
mysql> EXPLAIN ANALYZE
-> SELECT wb.brand_name, ROUND(AVG(w.retail_price), 2) AS avg_retail_price
-> FROM Watches w
-> IGNORE INDEX (PRIMARY, idx_retail_price_brand, idx_retail_price_brand, idx_brand_price, idx_brand_name_watches, idx_case_material, idx_brand_case) -- Ignore indexes on the Watches table
-> JOIN Watch_Brands wb
-> IGNORE INDEX (PRIMARY, idx_bq_country, idx_watch_brands_country, idx_brand_name_watch_brands, idx_group_brand) -- Ignore indexes on the Watch_Brands table
-> ON w.brand_name = wb.brand_name
-> GROUP BY wb.brand_name
-> HAVING AVG(w.retail_price) > (
->   SELECT AVG(bmp.Price)
->   FROM Brand_Market_Data bmp
->   WHERE bmp.brand_name = wb.brand_name
-> );
+-----+
| EXPLAIN
+-----+
|
+-----+
| -> Filter: (??? > '(select #2)') (actual time=107..107 rows=24 loops=1)
-> Table scan on <temporary> (actual time=105..105 rows=43 loops=1)
-> Aggregate using temporary table (actual time=105..105 rows=43 loops=1)
-> Inner hash join (wb.brand_name = w.brand_name) (cost=88715 rows=17384) (actual time=62.5..80.1 rows=17884 loops=1)
-> Table scan on wb (cost=217e-6 rows=50) (actual time=0.112..0.245 rows=50 loops=1)
-> Hash
-> Table scan on w (cost=1792 rows=17384) (actual time=0.101..9.05 rows=17884 loops=1)
-> Select #2 (subquery in projection; dependent)
-> Aggregate: avg(bmp.price) (cost=51.1 rows=1) (actual time=0.217..0.217 rows=1 loops=43)
-> Covering index lookup on bmp using idx_brand_market_data_brand_price (brand_name=wb.brand_name) (cost=28.7 rows=223) (actual time=0.103..0.195 rows=226 loops=43)
|
```

```
mysql> explain analyze SELECT wb.brand_name, ROUND(AVG(w.retail_price), 2) AS avg_retail_price
-> FROM Watches w
-> JOIN Watch_Brands wb ON w.brand_name = wb.brand_name
-> GROUP BY wb.brand_name
-> HAVING AVG(w.retail_price) > (
->   SELECT AVG(bmp.Price)
->   FROM Brand_Market_Data bmp
->   WHERE bmp.brand_name = wb.brand_name
-> );
+-----+
| EXPLAIN
+-----+
|
+-----+
| -> Filter: (avg(w.retail_price) > (select #2)) (cost=4887 rows=50) (actual time=1..24 rows=24 loops=1)
-> Group aggregate: avg(w.retail_price), avg(w.retail_price) (cost=4887 rows=50) (actual time=0.748..17.6 rows=43 loops=1)
-> Nested loop inner join (cost=2866 rows=20214) (actual time=0.589..12.1 rows=17884 loops=1)
-> Covering index scan on wb using idx_brand_name_watch_brands (cost=5.25 rows=50) (actual time=0.0491..0.101 rows=50 loops=1)
-> Covering index lookup on w using idx_retail_price_brand (brand_name=wb.brand_name) (cost=17.6 rows=404) (actual time=0.0357..0.215 rows=358 loops=50)
-> Select #2 (subquery in condition; dependent)
-> Aggregate: avg(bmp.price) (cost=51.1 rows=1) (actual time=0.145..0.145 rows=1 loops=43)
-> Covering index lookup on bmp using idx_brand_market_data_brand_price (brand_name=wb.brand_name) (cost=28.7 rows=223) (actual time=0.0456..0.123 rows=226 loops=43)
-> Select #2 (subquery in projection; dependent)
-> Aggregate: avg(bmp.price) (cost=51.1 rows=1) (actual time=0.145..0.145 rows=1 loops=43)
-> Covering index lookup on bmp using idx_brand_market_data_brand_price (brand_name=wb.brand_name) (cost=28.7 rows=223) (actual time=0.0456..0.123 rows=226 loops=43)
|
```

Query 4: Brands with >3 Case Material Variants

SQL Query

```
SELECT w.brand_name, COUNT(DISTINCT w.case_material) AS case_material_variants
FROM Watches w
JOIN Brand_Market_Data bmd ON w.brand_name = bmd.brand_name
GROUP BY w.brand_name
HAVING COUNT(DISTINCT w.case_material) > 3;
```

Indexes Tried

```
CREATE INDEX idx_brand_name_bmd ON Brand_Market_Data(brand_name);  
CREATE INDEX idx_case_material ON Watches(case_material);  
CREATE INDEX idx_brand_case ON Watches(brand_name, case_material);
```

Summary Table

Index Strategy	Cost Estimate	Actual Time (ms)
No indexes	17.1e+6	~6221
Index on brand_name (BMD)	886269	~5858
Index on case_material	886269	~5513
Index on (brand, material)	886269	~6007

Analysis

The cost estimate decreased significantly with the index on brand_name while remaining constant on case_material and a composite index.

This was by far our most intensive and taxing compute so it's good to see a large improvement.

```
mysql> EXPLAIN ANALYZE  
-> SELECT w.brand_name, COUNT(DISTINCT w.case_material) AS case_material_variants  
-> FROM Watches w  
-> IGNORE INDEX (PRIMARY, idx_retail_price_brand, idx_retail_price_brand, idx_brand_price, idx_brand_name_watches, idx_case_material, idx_brand_case) -- Ignore indexes on the Watches table  
-> JOIN Brand_Market_Data bmd  
-> IGNORE INDEX (PRIMARY, idx_brand_market_data_brand_price, idx_brand_name_bmd) -- Ignore indexes on the Brand_Market_Data table  
-> ON w.brand_name = bmd.brand_name  
-> GROUP BY w.brand_name  
-> HAVING COUNT(DISTINCT w.case_material) > 3;  
  
+-----+  
+-----+  
+-----+  
| EXPLAIN  
+-----+  
+-----+  
+-----+  
+-----+  
+-----+  
| -> Filter: (count(distinct Watches.case_material) > 3) (actual time=24809..29508 rows=37 loops=1)  
-> Group aggregate: count(distinct Watches.case_material), count(distinct Watches.case_material) (actual time=24808..29507 rows=43 loops=1)  
-> Sort: w.brand_name (actual time=24760..25389 rows=4.35e+6 loops=1)  
-> Stream results (cost=17.1e+6 rows=3.88e+6) (actual time=14.5..2356 rows=4.35e+6 loops=1)  
-> Inner hash join (bmd.brand_name = w.brand_name) (cost=17.1e+6 rows=3.88e+6) (actual time=14.5..590 rows=4.35e+6 loops=1)  
-> Table scan on bmd (cost=0.0166 rows=9833) (actual time=0.101..17.3 rows=9922 loops=1)  
-> Hash  
-> Table scan on w (cost=1792 rows=17384) (actual time=0.142..7.1 rows=17884 loops=1)  
|
```

```
mysql>
mysql> explain analyze SELECT w.brand_name, COUNT(DISTINCT w.case_material) AS case_material_variants
-> FROM Watches w
-> JOIN Brand_Market_Data bmd ON w.brand_name = bmd.brand_name
-> GROUP BY w.brand_name
-> HAVING COUNT(DISTINCT w.case_material) > 3;
+-----+
| EXPLAIN
+-----+
| -> Filter: (count(distinct w.case_material) > 3) (cost=886269 rows=43) (actual time=63.5..6063 rows=37 loops=1)
   -> Group aggregate: count(distinct w.case_material), count(distinct w.case_material) (cost=886269 rows=43) (actual time=62.7..6062 rows=43 loops=1)
     -> Nested loop inner join (cost=497776 rows=3.88e+6) (actual time=3.75..2331 rows=4.35e+6 loops=1)
       -> Covering index scan on w using idx_brand_case (cost=1792 rows=17384) (actual time=2.2..29.2 rows=17884 loops=1)
         -> Covering index lookup on bmd using idx_brand_name_bmd (brand_name=w.brand_name) (cost=6.18 rows=223) (actual time=0.0364..0.112 rows=243 loops=17884)
     |
+-----+
1 row in set (6.09 sec)
```